

Logistic regression

slides : https://docs.google.com/presentation/d/1OrjFYkvCl6HmYa8AS_glq_VRQXwG0BExa-fDr-I2MAE/edit#slide=id.p34

Your turn: Build a logistic regression model on the [penguins dataset] (<https://github.com/allisonhorst/palmerpenguins/tree/main?tab=readme-ov-file>) see also the Kaggle page for the dataset [Penguins on Kaggle] (<https://www.kaggle.com/datasets/parulpandey/palmer-archipelago-antarctica-penguin-data>) - load the data into a pandas dataframe - transform the categorical values into numerical ones - use a one vs rest or a one vs one strategy - using the best model - select the best regressors, are there any you can drop - assess the model performance - draw the distribution of probabilities for each class and assess the performance of the model - what samples are not properly classified - is logistic regression a good model ?
